

my_analytics

December 16, 2021

References: [Loan Approval Prediction](#)

Step 1: Set state number

```
[ ]: # replace 888 with the last three digits of your student id
# and then press CTRL + Enter
my_state_number = 673
```

Step 2: Read the data_set.csv

```
[ ]: import pandas as pd
data = pd.read_csv("data_set.csv")

data.head()
```

```
[ ]:      Loan_ID Gender Married Dependents      Education Self_Employed \
0  LP001002   Male      No           0      Graduate           No
1  LP001003   Male     Yes           1      Graduate           No
2  LP001005   Male     Yes           0      Graduate           Yes
3  LP001006   Male     Yes           0  Not Graduate           No
4  LP001008   Male      No           0      Graduate           No

      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
0                5849                0.0          NaN          360.0
1                4583             1508.0          128.0          360.0
2                3000                0.0           66.0          360.0
3                2583             2358.0          120.0          360.0
4                6000                0.0          141.0          360.0

      Credit_History  Property_Area  Loan_Status
0                1.0          Urban            Y
1                1.0          Rural            N
2                1.0          Urban            Y
3                1.0          Urban            Y
4                1.0          Urban            Y
```

Step 3: Sample the data randomly the data and save the dataframe as myNewData

```
[ ]: myNewData = data.sample(frac =.90, replace = False, random_state =
    ↳my_state_number)
```

Start the Analytics using **myNewData** dataframe as the raw data note: your *myNewData* dataframe may be different from other students' *myNewData* dataframe

1 0.0 Import dependencies

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
```

2 1.0 Data Exploration

2.1 1.1 Dataset Details

2.1.1 1.1.1 Data Shape

```
[ ]: # Start your codes
myNewData.shape
```

```
[ ]: (553, 13)
```

The data has 553 rows (data) and 13 columns (features).

2.1.2 1.1.2 Data Information

```
[ ]: myNewData.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 553 entries, 461 to 66
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                553 non-null   object
1   Gender                 541 non-null   object
2   Married                551 non-null   object
3   Dependents             540 non-null   object
4   Education              553 non-null   object
5   Self_Employed          523 non-null   object
6   ApplicantIncome        553 non-null   int64
```

```

7   CoapplicantIncome  553 non-null    float64
8   LoanAmount         535 non-null    float64
9   Loan_Amount_Term   541 non-null    float64
10  Credit_History     507 non-null    float64
11  Property_Area      553 non-null    object
12  Loan_Status        553 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 60.5+ KB

```

All columns are of the object datatype except for ApplicantIncome that has int64 type, CoapplicantIncome, LoanAmount, Loan_Amount_Term, Credit_History of float64 type. There are missing data where some columns do not have all 553 rows. The column names are inconsistent where some uses underscores and some using pure CamelCase.

2.1.3 1.1.3 Index Dropping

Dropping Loan_ID before any data analysis is conducted as it does not offer any real meaning.

```
[ ]: myNewData = myNewData.drop(columns='Loan_ID')
```

2.1.4 1.1.4 Column Renaming

Renaming of data columns for a more consistent experience. CamelCase will be used.

```
[ ]: myNewData = myNewData.rename(columns={'Self_Employed': 'SelfEmployed',
↪ 'Loan_Amount_Term': 'LoanAmountTerm', 'Credit_History': 'CreditHistory',
↪ 'Property_Area': 'PropertyArea', 'Loan_Status': 'LoanStatus'})
```

2.1.5 1.1.5 Nature of Data

```
[ ]: myNewData.head()
```

```
[ ]:
      Gender Married Dependents      Education SelfEmployed  ApplicantIncome \
461   Male      Yes          3+      Graduate             No             7740
597   Male      No          NaN      Graduate             No             2987
455   Male      Yes          2      Graduate             No             3859
6     Male      Yes          0   Not Graduate             No             2333
196   Male      No          0      Graduate             No             8333
```

```

      CoapplicantIncome  LoanAmount  LoanAmountTerm  CreditHistory \
461                  0.0        128.0          180.0           1.0
597                  0.0         88.0          360.0           0.0
455                  0.0         96.0          360.0           1.0
6                   1516.0         95.0          360.0           1.0
196                 3750.0        187.0          360.0           1.0

```

```

      PropertyArea LoanStatus
461          Urban          Y

```

597	Semiurban	N
455	Semiurban	Y
6	Urban	Y
196	Rural	Y

From the table above, it is known that numerical variables are `ApplicantIncome`, `CoapplicantIncome` and `LoanAmount`. The rest are categorical variables.

```
[ ]: cat_data = (myNewData.drop(columns=['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount'])).columns.values
num_data = (myNewData[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount']]).columns.values
```

2.1.6 1.1.6 Data Description

```
[ ]: myNewData.describe()
```

```
[ ]:
ApplicantIncome  CoapplicantIncome  LoanAmount  LoanAmountTerm \
count      553.000000      553.000000  535.000000      541.000000
mean      5479.320072     1577.437468   147.319626     341.656192
std       6357.052607     2916.184573    87.882458     65.922795
min       150.000000       0.000000     9.000000     12.000000
25%      2833.000000       0.000000   100.000000     360.000000
50%      3813.000000     1086.000000   128.000000     360.000000
75%      5780.000000     2250.000000   167.500000     360.000000
max      81000.000000    41667.000000   700.000000     480.000000

CreditHistory
count      507.000000
mean        0.852071
std         0.355380
min         0.000000
25%         1.000000
50%         1.000000
75%         1.000000
max         1.000000
```

The statistics above describe numerical data columns. It can be known that the data values of these columns are spread across a large range. For example, `ApplicantIncome` has a maximum value of 81000 whereas `CreditHistory` has a maximum value of only 1. These data values need to be scaled for better model accuracy. Algorithms based on gradient descent such as linear regression and neural network perform better with scaled data because the data values will affect the step size of the gradient descent. The gradient descent will converge more quickly towards the minima when using data on a similar scale. Distance-based algorithms such as KNN and SVM are most affected by the range of data values as they calculate the distances between data points to find the similarity. The algorithms will stress more on features with data of a higher value, causing the model to be biased. Tree-based algorithms are quite insensitive to the data scales because the tree splits on a

feature without taking other features into consideration. There are two main scaling techniques, that are normalisation and standardisation. Normalisation will transform all values to fit in the range of 0 and 1, also known as min-max scaling. Standardisation turns the mean value into 0 and the other values centred around the mean value will have a unit standard deviation. There is no particular range to this scaling method. Normalisation is used when the data distribution does not follow a Gaussian distribution, especially for KNN and neural networks, but it is very prone to outliers. Standardisation is helpful when the data follows a Gaussian distribution, but it is not necessarily so. Outliers in the data will not be affected by standardisation. The mean and standard deviation will be rescaled in such a way that they are very close to 0 and 1 respectively.

2.2 1.2 Data Exploration

2.2.1 1.2.1 Data Distribution

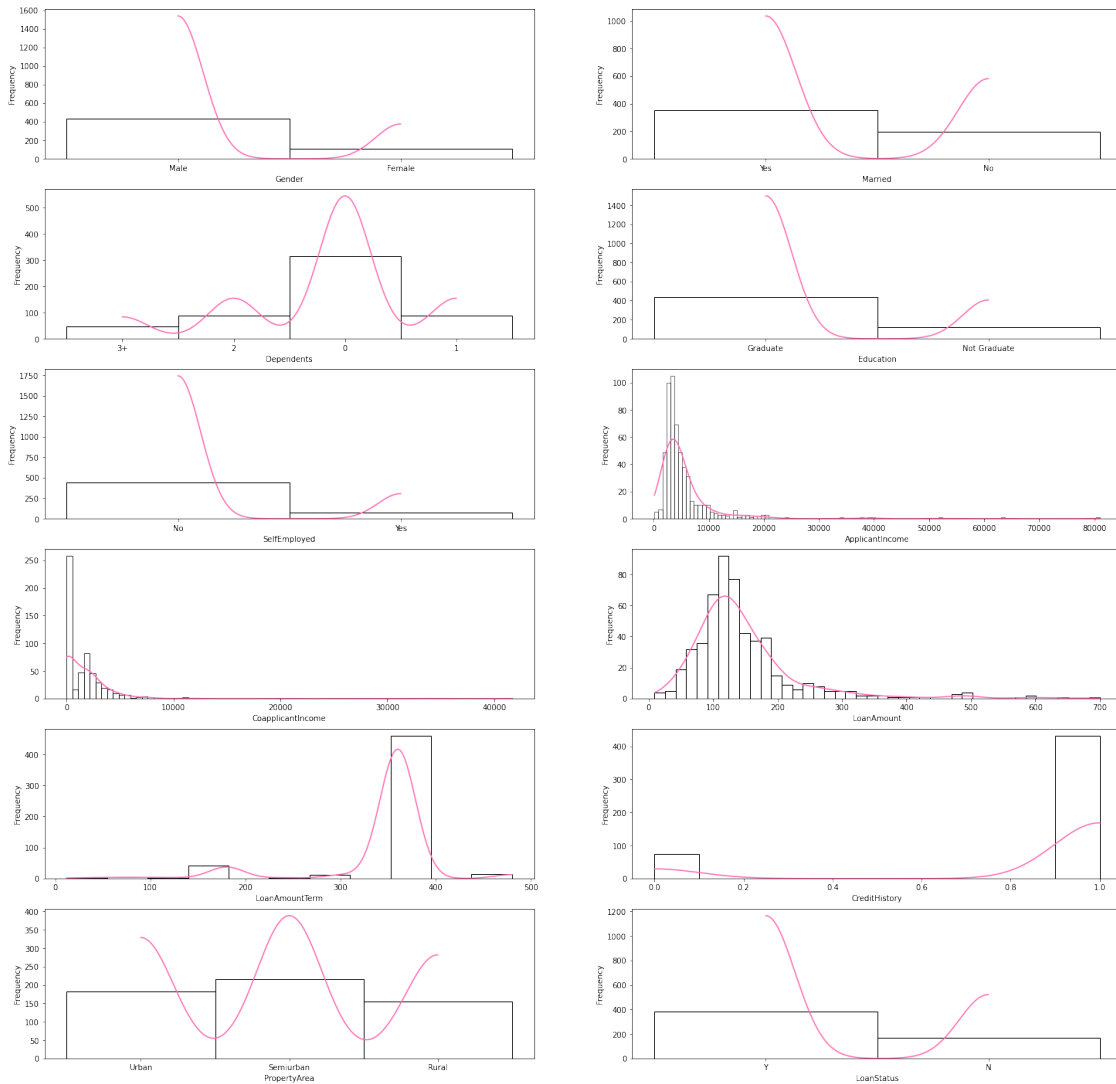
Exploring the dataset's general distribution pattern.

```
[ ]: cols = list(myNewData.columns.values)

fig, ax = plt.subplots(6, 2, figsize=(25,25))
# fig.suptitle('Distribution of Dataset')

ax = ax.flatten() # ax is flattened from a 2D array to a 1D array, use ax.T.
    ↳ flatten() to transpose if needed

for i in range(len(cols)):
    sns.histplot(data=myNewData, ax=ax[i], x=cols[i], kde=True,
    ↳ color='hotpink', alpha=0)
    ax[i].set(xlabel=cols[i], ylabel='Frequency')
```



As shown from the graph above, there are five data distributions that do not follow a bell-shaped curve, which are all binary variables. All other data follows a normal distribution pattern. This analysis suggests that data standardisation instead of data normalisation might be applied onto the dataset.

2.2.2 1.2.1 Categorical Data Analysis

Analysing categorical variables. Exploring the relationships between the features and the target variable (LoanStatus).

```
[ ]: fig, ax = plt.subplots(4, 2, figsize=(25,25))

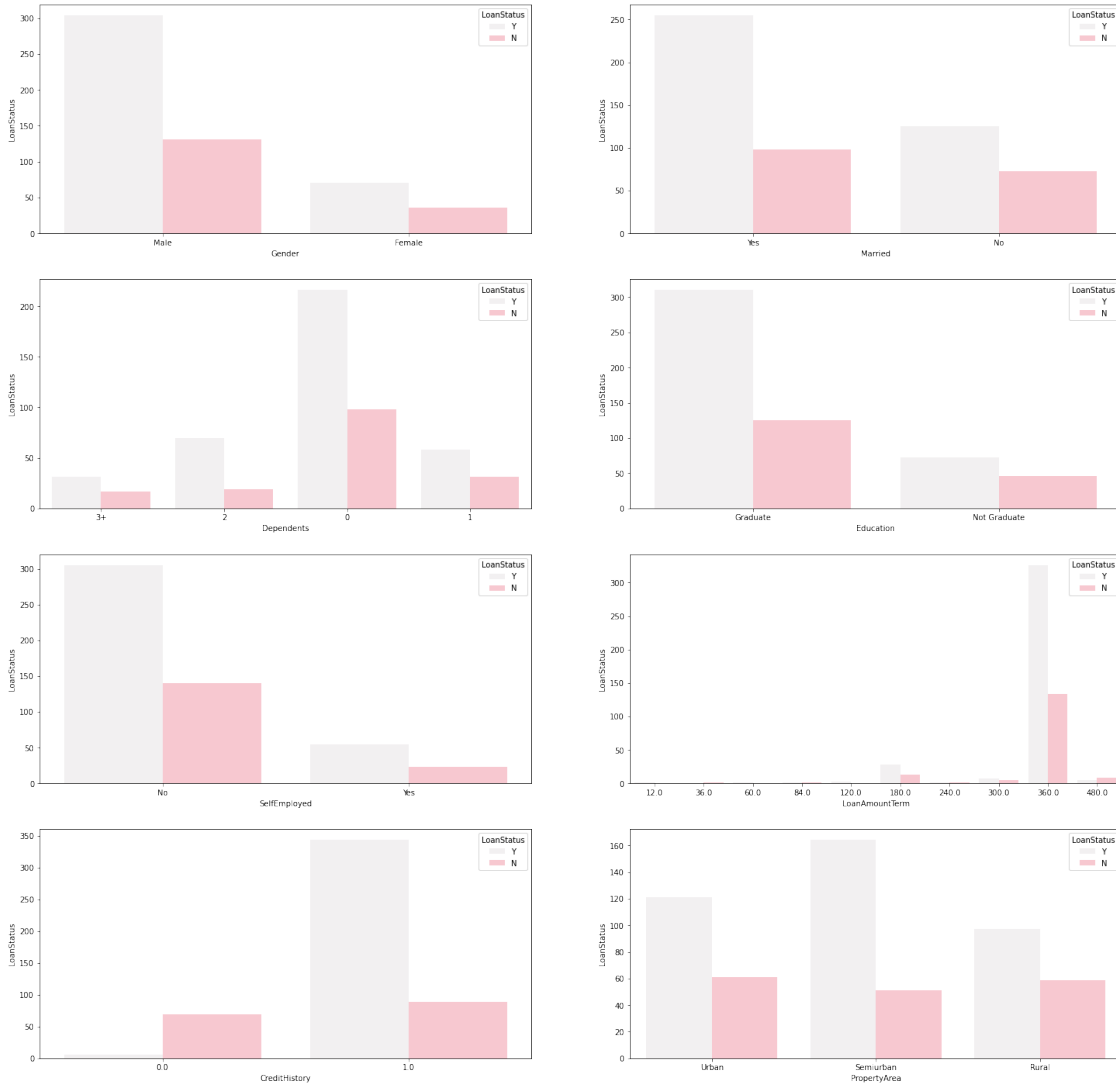
ax = ax.flatten()

for i, v in enumerate(cat_data):
```

```

if(v == 'LoanStatus'):
    continue
sns.countplot(data=myNewData, ax=ax[i], x=v, hue='LoanStatus',
color='pink', alpha=1)
ax[i].set(xlabel=v, ylabel='LoanStatus')

```



1. Around 3/5 applications have been approved.
2. There are approximately 3 times more male applicants than female.
3. Around 3/5 applicants are married. Married applicants are more likely to be granted loans.
4. Around 3/5 of the applicants have zero dependents. They are more likely to be granted loans.
5. Around 4/5 of the applicants are graduates, who are also more likely to be granted loans.
6. Less than 1/5 of the applicants are self-employed, who are not as likely to be granted loans as their counterparts.
7. Majority of the applicants have applied for a 30-year loan (360 months).

8. Applicants without credit history are unlikely to be granted loans.
9. There are more applicants from the semiurban property area and they are more likely to be granted loans.

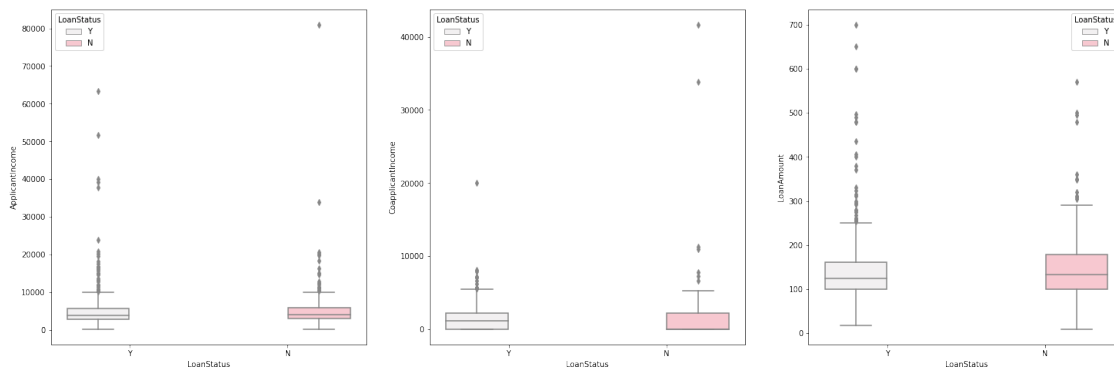
2.2.3 1.2.2 Numerical Data Analysis

Analysing numerical columns.

```
[ ]: fig, ax = plt.subplots(1, 3, figsize=(25,8))

ax = ax.flatten()

for i, v in enumerate(num_data):
    if(v == 'LoanStatus'):
        continue
    sns.boxplot(data=myNewData, ax=ax[i], x='LoanStatus', y=v,
        hue='LoanStatus', color='pink')
    ax[i].set(xlabel='LoanStatus', ylabel=v)
```



It seems that these numerical variables do not have a significant relationship to **LoanStatus** as the boxes of Y and N do not have a significant difference between them.

3 2.0 Data Preprocessing

3.0.1 Null Value Check

Checking for null values.

3.0.2 Data Encoding

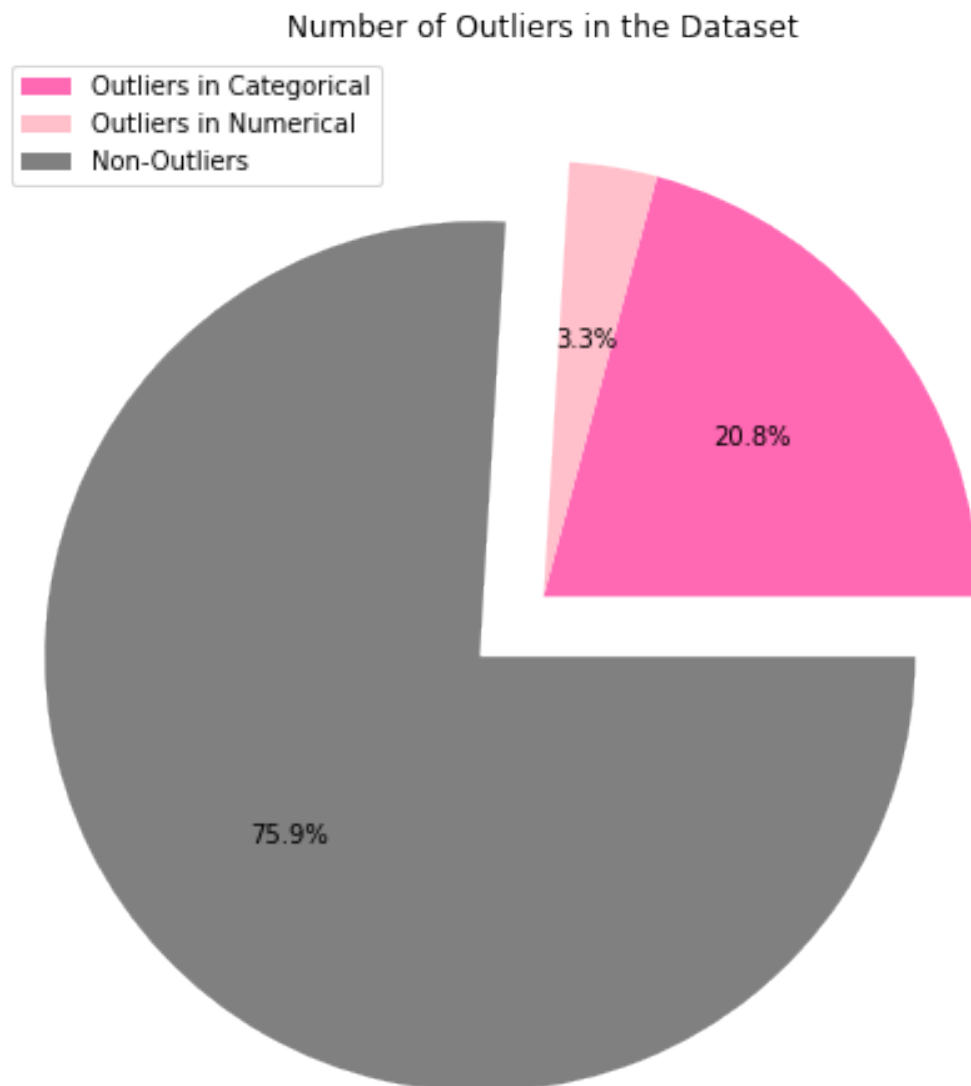
Encoding categorical data values to be numerical for easier further processing. Label encoding is used for binary categorical columns such as **gender**, and these feature columns are renamed in such a way that 0 is false and 1 is true for the column to replicate the effect of one hot encoding. **PropertyArea** and **Dependents** columns will be encoded using one-hot encoding.


```
[ ]: binary = ['Gender', 'Married', 'Education', 'SelfEmployed', 'CreditHistory',
↳ 'LoanStatus']

label_encoder = LabelEncoder()

for _, v in enumerate(binary):
    myNewData[v] = label_encoder.fit_transform(myNewData[v])

[ ]: myNewData2 = myNewData.rename(columns={'Gender': 'Male', 'Education':
↳ 'Graduate', 'LoanStatus': 'LoanApproved'})
myNewData2
```



```
[ ]: onehot = ['Dependents', 'PropertyArea']

myNewData3 = myNewData2

onehot_encoder = OneHotEncoder(sparse=False, drop=None)

for _, v in enumerate(onehot):
    tmp_df = pd.DataFrame(onehot_encoder.fit_transform(np.reshape(myNewData2[v].
    ↳values, (-1, 1))))
    myNewData3 = myNewData3.drop(columns=[v])
    myNewData3 = myNewData3.join(tmp_df, rsuffix='_' + v)

myNewData3
```

Outliers in Categorical: 115
 Outliers in Numerical: 18
 Total Outliers: 133

```
[ ]: df_encoded = myNewData.replace({'Male': 0, 'Female': 0, 'No': 0, 'Yes': 1, 'Not_
    ↳Graduate': 0, 'Graduate': 1, 'Rural': 0, 'Semiurban': 1, 'Urban': 2, 'N': 0,
    ↳'Y': 1, '0': 0, '1': 1, '2': 2, '3+': 3})
```

3.0.3 Simple Imputation

There are a total of 115 rows of data which is 20.8% percent of the data that are categorical in nature. These rows will be imputed using the `SimpleImputer` class by using the data with the highest frequency. On the other hand, numerical data will be imputed with the mean of the column.

```
[ ]: for _, v in enumerate(cat_data):
    myNewData[v] = myNewData[v].fillna(myNewData[v].value_counts().index[0]) #
    ↳Fill categorical data with the highest frequency

for _, v in enumerate(num_data):
    myNewData[v] = myNewData[v].fillna(myNewData[v].mean()) # Fill numerical
    ↳data with column mean

myNewData.isnull().sum()
```

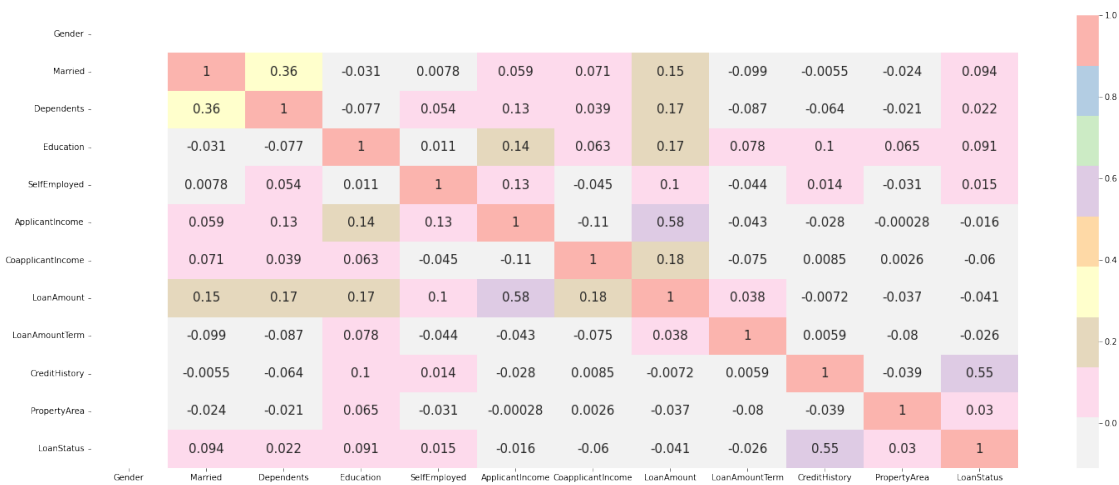
```
[ ]: Gender          0
     Married         0
     Dependents      0
     Education       0
     SelfEmployed    0
     ApplicantIncome 0
     CoapplicantIncome 0
     LoanAmount      0
```

```
LoanAmountTerm      0
CreditHistory        0
PropertyArea         0
LoanStatus           0
dtype: int64
```

Heatmap

```
[ ]: plt.figure(figsize=(25, 10))
sns.heatmap(df_encoded.corr(), cmap='Pastel1_r', annot=True, annot_kws={'size': 15})
```

[]: <AxesSubplot:>



```
[ ]: myNewData.isnull().sum()
```

```
[ ]:
Male    Married  Dependents  Graduate  SelfEmployed  ApplicantIncome  \
461      1         1          3+           0             0             7740
597      1         0          0           0             0             2987
455      1         1          2           0             0             3859
6         1         1          0           1             0             2333
196      1         0          0           0             0             8333
..      ...      ...      ...      ...      ...      ...
227      1         1          2           0             0             6250
566      1         0          0           0             0             3333
229      1         0          0           0             1             6400
408      1         1          1           0             0             8300
66       1         0          0           1             0             3200

CoapplicantIncome  LoanAmount  LoanAmountTerm  CreditHistory  \
461                0.0        128.0            180.0            1
```

597	0.0	88.0	360.0	0
455	0.0	96.0	360.0	1
6	1516.0	95.0	360.0	1
196	3750.0	187.0	360.0	1
..
227	1695.0	210.0	360.0	1
566	0.0	70.0	360.0	1
229	0.0	200.0	360.0	1
408	0.0	152.0	300.0	0
66	2254.0	126.0	180.0	0

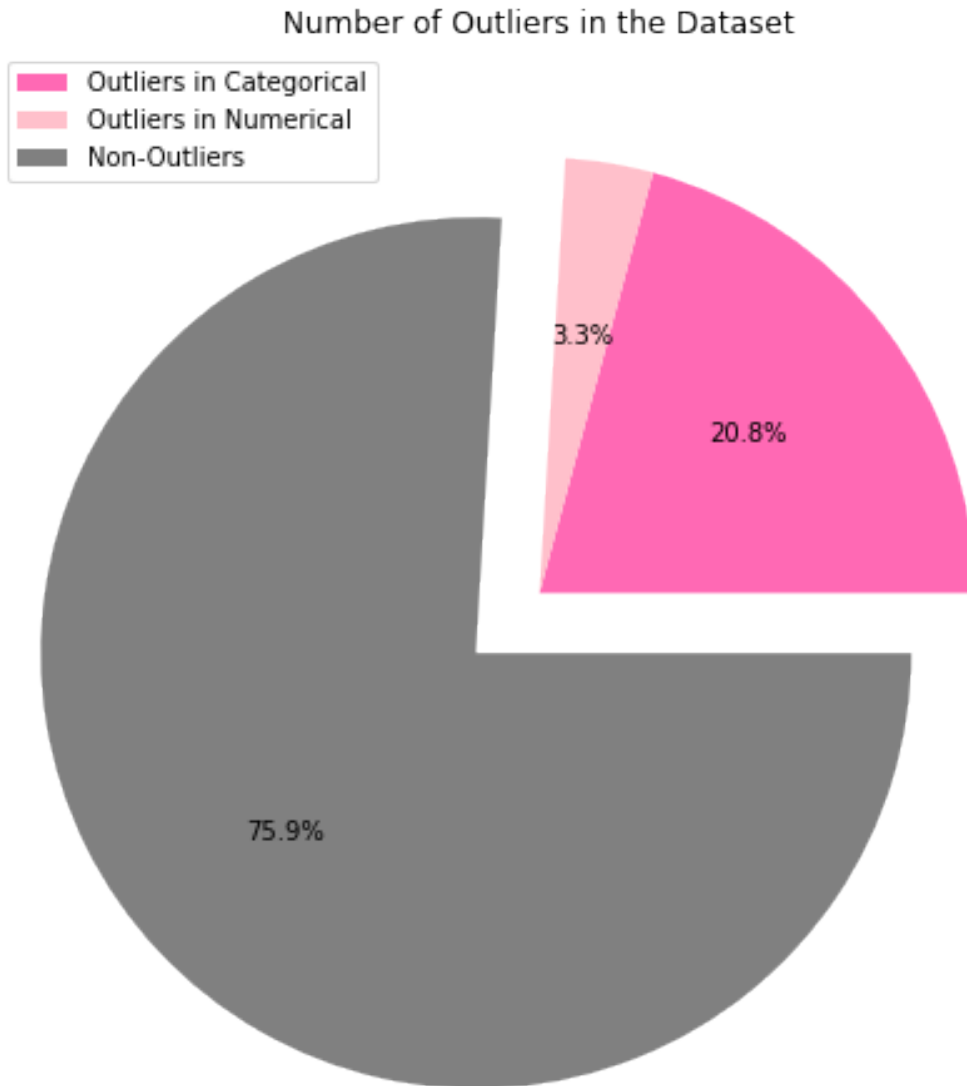
	PropertyArea	LoanApproved
461	Urban	1
597	Semiurban	0
455	Semiurban	1
6	Urban	1
196	Rural	1
..
227	Semiurban	1
566	Urban	1
229	Rural	1
408	Semiurban	0
66	Urban	0

[553 rows x 12 columns]

There are null values in the columns Gender, Married, Dependents, SelfEmployed, LoanAmount, LoanAmountTerm, CreditHistory.

```
[ ]: x = ['Outliers in Categorical', 'Outliers in Numerical', 'Non-Outliers']
y = [myNewData[cat_data].isnull().values.sum(), myNewData[num_data].isnull().
     ↪ values.sum(), myNewData.shape[0] - myNewData.isnull().values.sum()]

fig = plt.gcf()
fig.set_size_inches(8, 8)
plt.pie(y, colors=['hotpink', 'pink', 'gray'], explode=[0, 0, 0.2], autopct='%1.
     ↪ 1f%%')
plt.legend(x)
plt.title("Number of Outliers in the Dataset")
plt.show()
```



```
[ ]: print('Outliers in Categorical:', myNewData[cat_data].isnull().values.sum())
      print('Outliers in Numerical:', myNewData[num_data].isnull().values.sum())
      print('Total Outliers:', myNewData.isnull().values.sum())
```

```
Outliers in Categorical: 115
Outliers in Numerical: 18
Total Outliers: 133
```

The total number of rows with missing data is 133, which is almost 1/5 of the dataset. Removing all these data might cause biasness in the model as the remaining dataset will be quite small. Therefore, the treatment approach will be taken.

4 2.0 Data Preprocessing

Dealing with missing values.

```
[ ]: df_filled = df_encoded.fillna(df_encoded.mean())
df_filled.isnull().any()
```

```
[ ]: Gender                False
Married                  False
Dependents               False
Education                False
SelfEmployed            False
ApplicantIncome          False
CoapplicantIncome        False
LoanAmount              False
LoanAmountTerm           False
CreditHistory           False
PropertyArea            False
LoanStatus              False
dtype: bool
```

The numerical columns are filled with the mean of the columns from the original data before dropping any gender columns. Using mean is due to that it best represents the average. Calculating mean from the original data before dropping rows is because those dropped data are still valid and valuable data that contributes meaningfully to the mean.

```
[ ]: # df = df_filled.drop()
df = df_filled
```

```
[ ]: X_df = df.drop('LoanStatus', axis=1)
y_df = df['LoanStatus']
```

Split features and target label.

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X_df, y_df, test_size = 0.
→3, random_state=0)
```

Split dataset into training and test sets using the 70:30 ratio for training:testing.

```
[ ]: scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Normalising the data to improve model accuracy.

5 Model Selection